

Data-driven *pricing*

Starting with luxury event venues, where every date is a one-off and every price is a negotiation.

01 THE PARTNER

A Manhattan event venue selling event days at

\$50K–\$200K

Galas, conferences and celebrations, booked one date at a time.

02 THE INCUMBENT

A static rate card

Varies only by year, month and weekday. It never sees:

Event type

Lead time

Demand pressure

Every quote is negotiated by hand. The hardest call, *take this deal or hold the date for a better one*, is made on gut feel.

03 MLIQ

Replace the card with models trained on the venue's own sales history.

Delivered inside the existing Salesforce workflow, so pricing and lead scoring live where the sellers already work.

Predict, then *prescribe*

Stage 1 predicts what shows up and what it's worth. Stage 2 prescribes the call.

PREDICTIVE · MODEL 1 · DEMAND

What kind of demand does this date attract?

Event type and calendar footprint. A conference books midweek and consumes 2–4 days with load-ins; a gala takes one.

PREDICTIVE · MODEL 2 · VALUE

What should this date cost?

By event type, section, weekday and month, as a **negotiable range**, not a single number.

PRESCRIPTIVE · PRICING THE OPPORTUNITY COST

Accepting a one-day social on a Wednesday destroys the **Wed+Thu** and **Tue+Wed** conference blocks, the two most common multi-day shapes in the book.

\$195K


MEDIAN BLOCK VALUE

ONE DECISION RULE

Hold the date if $P(\text{better arrives}) \times E[\text{value}] - \text{holding cost} > \text{the offer in hand}$

THE BAR TO CLEAR · MAE, LOWER IS BETTER

56.4% HEAD-TO-HEAD VS THE RATE CARD



Curves fitted to realized bookings already beat the card, but only *in-sample today*. The bake-off re-measures it on an out-of-time holdout. **The models must clear the fitted curves, not the card.**

One venue's history, *2021–2026*

Extracted and anonymized end-to-end: pseudonyms, an allow list of approved fields, and no names, emails or free text leaving the CRM.

3,985

OPPORTUNITIES • THE DEALS

5,877

RESERVATIONS • THE CALENDAR

1,170

LOST OPPORTUNITIES RESTORED

\$131M

QUOTED-AND-LOST REVENUE RECOVERED IN CLEANUP

THE HOLE

2021–2023 was migrated from a legacy CRM *as wins only*.
The lost deals existed as empty shells, or not at all.

266


REBUILT FROM SCRATCH

904

EMPTY SHELLS REPAIRED

USABLE DEMAND HISTORY

BEFORE • AS DELIVERED



AFTER • RECONSTRUCTED



2021 2022 2023 2024 2025 2026

Demand history now runs from 2022, not 2024. The modeling window doubled.

Finding the *grain*

The unit of analysis had to be earned. The calendar lies about money.

THE TRAP

Reservation money fields are formulas that read the **parent Opportunity**. Summing day money off the calendar inflates revenue

2.5×

THE FIX

Rebuild per-day dollars from the **product lines underneath**, which reconcile to opportunity totals to the cent.

Reconciled to the cent

THE TARGET

Venue rent per event day

\$80,000

MEDIAN WON · FULL VENUE \$90K · PARTIAL \$34.75K

MODELING GRAIN 1

3,706

EVENTS · ONE PER INQUIRY: OUTCOME, EVENT TYPE, TRUE INQUIRY AND LOSS DATES

MODELING GRAIN 2

5,851

EVENT-DAYS · ONE PER DAY PER VENUE SECTION: ITS OWN HOLD POSITION, ITS OWN DOLLARS

OUTCOMES

659 2,981 66

WON LOST OPEN




4.5:1 lost to won. Exactly why a model that learns only from wins **misreads** the market.

Inquiry to *Booked*


Every inquiry lands as a Lead. Once qualified it becomes an Opportunity that holds dates on the master calendar.

01 CAPTURE • ANY CHANNEL, ONE ENTRY POINT



Each Lead Reservation looks up **one** Lead Product, and the Lead.

02 QUALIFY • SALES REVIEWS EVERY LEAD



Pipeline, holds and *money*

The money rides on the products attached to each day.

03 PIPELINE • OPPORTUNITY STAGES

STAGE 1 Qualified	STAGE 2 Proposal	STAGE 3 Negotiation	STAGE 4 Contract	WON HERE Booked	DELIVERED Closed Won
----------------------	---------------------	------------------------	---------------------	--------------------	-------------------------

Won at Booked. Contract signed, deposit in. Closed Won is bookkeeping after delivery and deposit refund.

Closed Lost. Exit from any stage, mostly on *date availability* and *price*.

04 CALENDAR AND MONEY

MASTER CALENDAR

Reservation

Each product on the Opportunity, every day, load-in included

CHILD OF OPPORTUNITY

Opportunity Product

Venue rent · Beverage · Staffing. One row per item, per day

Per-day pricing. Money sits on each day's line items, not the Opportunity as a whole.

HOLD POSITION • ONE DATE, MANY DEALS

1st First claim on the date

2nd Moves up if the 1st is released

3rd Still tracked and quotable

WIN

Holds flip to **Confirmed**; date blocked.

LOSS

Date released; **next hold moves up**.

The two *predictive* models

Different grains, different training sets.
One decision.

	MODEL 1	MODEL 2
	Demand	Value
ESTIMATES	Arrival rate λ per demand shape, by weekday, month and days-out. P(better deal in k days)	Venue rent per event day, as a low / point / high range, because every price here is negotiated.
ALGORITHM	Poisson arrival model, partial pooling across sparse shape cells	Gradient-boosted quantile regression on log rent/day
GRAIN	Event (3,706) . An event arrives once, then occupies days; day grain would count a 3-day conference three times	Event-day (5,851) . The same event's Thursday and Saturday are worth different money
FEATURES	weekday, month, lead time, demand shape (type $\times$ footprint), market segment	event type, venue section, weekday, month, guest count, lead time, footprint length, segment, industry
TRAINS ON	All arrivals, won AND lost . A deal lost to a blocked date is displaced demand, not a price rejection	Priced won days + clean rejections only; blocked-date losses excluded, because their high quotes would teach “ <i>high prices lose</i> ”

THE DEMAND SHAPES ARE REAL CELLS, NOT THEORY

1-day conference-family		973
1-day gala		583
1-day reception		452
1-day wedding		362
2-day conference		244
3-day conference		73

Thin tails share strength through partial pooling; the **multi-day corporate shapes** are the \$195K blocks the opportunity-cost math protects.

ONE SHARED TRAIN / SERVE CONTRACT

9 required + 5 optional features, with leakage forbidden by construction: rules-engine outputs, post-outcome fields and dead source fields **raise an error** if they reach a model.

Alternatives

Every candidate is judged on the same rows, the same splits, once.

ALTERNATIVE	GRAIN	VERDICT
Multinomial event-type model Advisor's suggestion: $P(\text{event type} \mid \text{weekday, month, year})$	Event	ADOPTED As the demand-shape baseline inside Model I, but arrival intensity stays separate: <i>what</i> books Wednesdays and <i>how often</i> anything arrives move independently.
Summed one-day opportunity costs Price a hold as the sum of individual day costs	Event-day	KEPT AS LOWER BOUND It undercounts block displacement, because a Wednesday booking kills every multi-day block crossing Wednesday, so destroyed blocks are priced jointly at event grain.
Neural networks Feed-forward net for value; neural point-process for arrivals	Same as the champion each challenges	RUN AS CHALLENGERS Viable now the reconstruction doubled the history. Identical rows and partitions, promoted only on shared-holdout wins.

ONE PARTITION DESIGN • OUT-OF-TIME 60 / 20 / 20



Ordered by event date, cut at event boundaries, grouped by account, so **no client and no multi-day event straddles a split**. The final 20% is one shared holdout, **scored once**.

WHY GBM STARTS CHAMPION

On ~3,700 tabular events: native quantile loss for negotiated ranges, monotonic constraints, and attributions a revenue manager can audit.

Whatever wins must also beat the re-fit curves on the holdout. **Nothing ships on architecture appeal.**

What we're claiming, and what would *prove us wrong*

The data work is done. What's left is the measurement we've already committed to.

01 THE CLAIM

A model trained on this venue's own history prices a date better than a card that cannot see event type, lead time or demand pressure.

And it prices the harder question the card never asks: *hold, or take the offer in hand?*

NEXT • 01

Score the out-of-time holdout **once**, and report whether the models clear the re-fit curves.

02 THE PROOF

One out-of-time holdout, scored once.

Champions must beat the **re-fit curves**, not the rate card. Beating the card is the floor, not the result.

No leakage by construction; nothing promoted on architecture appeal.

NEXT • 02

Run the challengers on identical rows and partitions. Promote only on a **shared-holdout win**.

03 THE FAILURE WE ALREADY PRICED

A model that learns only from wins.

4.5:1 lost to won is why the 1,170 lost opportunities were rebuilt *before* anything was fit.

A blocked-date loss is displaced demand, not a price rejection. The two models split on exactly that line.

NEXT • 03

Put the quote and the hold call in front of sellers **inside Salesforce**, where the work already happens.